

```
In [3]: # Importing important libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime as dt
import xverse
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from dmba import classificationSummary, regressionSummary
from dmba import classificationSummary, gainsChart, liftChart
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, classification_report
%matplotlib inline
```

no display found. Using non-interactive Agg backend

```
In [4]: org=pd.read_csv('org.csv', parse_dates=["date_of_joining", "last_working_date", "cutoff_date", "date_of_birth"])
org.head(5)
```

Out[4]:

	emp_id	status	turnover	location	level	date_of_joining	date_of_birth	last_working_date
0	E11061	Inactive	1	New York	Analyst	2012-03-22	1992-03-22	2014-11-09
1	E1031	Inactive	1	New York	Analyst	2012-09-03	1992-10-01	2014-05-06
2	E6213	Inactive	1	New York	Analyst	2012-06-01	1992-06-02	2014-04-30
3	E5900	Inactive	1	New York	Analyst	2012-03-22	1991-12-19	2014-09-04
4	E3044	Inactive	1	Florida	Analyst	2012-03-29	1991-10-12	2014-01-23

```
In [5]: org.shape
```

Out[5]: (2291, 14)

```
In [6]: org.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2291 entries, 0 to 2290
Data columns (total 14 columns):
emp_id                2291 non-null object
status                2291 non-null object
turnover              2291 non-null int64
location              2291 non-null object
level                 2291 non-null object
date_of_joining       2291 non-null datetime64[ns]
date_of_birth         2291 non-null datetime64[ns]
last_working_date     410 non-null datetime64[ns]
gender                2291 non-null object
department            2291 non-null object
mgr_id                2291 non-null object
cutoff_date           2291 non-null datetime64[ns]
generation            2291 non-null object
emp_age               2291 non-null float64
dtypes: datetime64[ns](4), float64(1), int64(1), object(8)
memory usage: 250.7+ KB
```

## Data cleaning

Changing data types of some columns \*\*Filling NaN values with 0

```
In [7]: #print(today_date)
org["last_working_date"] = org["last_working_date"].fillna(pd.to_datetime(
("2021-03-12")))
org.head(10)
```

Out[7]:

	emp_id	status	turnover	location	level	date_of_joining	date_of_birth	last_working_date
0	E11061	Inactive	1	New York	Analyst	2012-03-22	1992-03-22	2014-11-09
1	E1031	Inactive	1	New York	Analyst	2012-09-03	1992-10-01	2014-05-06
2	E6213	Inactive	1	New York	Analyst	2012-06-01	1992-06-02	2014-04-30
3	E5900	Inactive	1	New York	Analyst	2012-03-22	1991-12-19	2014-09-04
4	E3044	Inactive	1	Florida	Analyst	2012-03-29	1991-10-12	2014-01-23
5	E4008	Active	0	Florida	Assistant Manager	2013-11-27	1992-06-05	2021-03-12
6	E6636	Active	0	New York	Specialist	2012-02-17	1992-01-23	2021-03-12
7	E13796	Inactive	1	New York	Analyst	2012-03-30	1990-12-19	2014-11-05
8	E13549	Active	0	New York	Analyst	2012-09-03	1991-12-22	2021-03-12
9	E13430	Inactive	1	New York	Analyst	2012-09-03	1991-08-19	2014-10-16

```
In [8]: org.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2291 entries, 0 to 2290
Data columns (total 14 columns):
emp_id          2291 non-null object
status          2291 non-null object
turnover        2291 non-null int64
location        2291 non-null object
level           2291 non-null object
date_of_joining 2291 non-null datetime64[ns]
date_of_birth   2291 non-null datetime64[ns]
last_working_date 2291 non-null datetime64[ns]
gender          2291 non-null object
department      2291 non-null object
mgr_id          2291 non-null object
cutoff_date     2291 non-null datetime64[ns]
generation      2291 non-null object
emp_age         2291 non-null float64
dtypes: datetime64[ns](4), float64(1), int64(1), object(8)
memory usage: 250.7+ KB
```

# Exploratory Data Analysis

## Time frame of data

```
In [9]: time_max=org["cutoff_date"].max()  
time_min=org["date_of_joining"].min()  
print("End time for data collection", time_max)  
print("Starting time for data collection", time_min)
```

```
End time for data collection 2014-12-31 00:00:00  
Starting time for data collection 1994-01-21 00:00:00
```

```
In [10]: # Number of employees at the start of the 20 year time period- active and inactive  
org["status"].value_counts()
```

```
Out[10]: Active      1881  
Inactive    410  
Name: status, dtype: int64
```

**Turnover rate= Total number of inactive employees / (Active employees at beginning of time period+Active employees at end of time period)/2**

```

In [11]: count_2004_active=org[(org["date_of_joining"]<="2004-12-31")&
                                     (org["last_working_date"]>"2004-12-31")][ "emp_id" ]
        .count()
count_2004_inactive=org[(org["date_of_joining"]<="2004-12-31")
                           &(org["last_working_date"]<"2004-12-31")][ "emp_id" ]
        .count()
print("Number of active employees in 2004: ",count_2004_active)
print("Number of inactive employees in 2004: ",count_2004_inactive)
print("\n")

count_2005_active = org[(org["date_of_joining"]<="2005-12-31") &
                           (org["last_working_date"]>="2005-12-31")][ "emp_id" ]
        .count()
count_2005_inactive = org[(org["date_of_joining"]<="2005-12-31") &
                           (org["last_working_date"]<="2005-12-31")][ "emp_id" ]
        .count()
print("Number of active employees in 2005: ",count_2005_active)
print("Number of inactive employees in 2005: ",count_2005_inactive)
print("\n")

count_2006_active=org[(org["date_of_joining"]<="2006-12-31") &
                           (org["last_working_date"]>="2006-12-31")][ "emp_id" ]
        .count()
count_2006_inactive=org[(org["date_of_joining"]<="2006-12-31") &
                           (org["last_working_date"]<="2006-12-31")][ "emp_id" ]
        .count()
print("Number of active employees in 2006: ",count_2006_active)
print("Number of inactive employees in 2006: ",count_2006_inactive)
print("\n")

count_2013_active=org[(org["date_of_joining"]<="2013-12-31") &
                           (org["last_working_date"]>="2013-12-31")][ "emp_id" ]
        .count()
count_2013_inactive=org[(org["date_of_joining"]<="2013-12-31") &
                           (org["last_working_date"]<="2013-12-31")][ "emp_id" ]
        .count()
print("Number of active employees in 2013: ",count_2013_active)
print("Number of inactive employees in 2013: ",count_2013_inactive)
print("\n")

count_2014_active=org[(org["date_of_joining"]<="2014-12-31") &
                           (org["last_working_date"]>="2014-12-31")][ "emp_id" ]
        .count()
count_2014_inactive=org[(org["date_of_joining"]<="2014-12-31") &
                           (org["last_working_date"]<="2014-12-31")][ "emp_id" ]
        .count()
print("Number of active employees in 2014: ",count_2014_active)
print("Number of inactive employees in 2014: ",count_2014_inactive)

```

Number of active employees in 2004: 148  
 Number of inactive employees in 2004: 0

Number of active employees in 2005: 206  
 Number of inactive employees in 2005: 0

Number of active employees in 2006: 329  
 Number of inactive employees in 2006: 0

Number of active employees in 2013: 2291  
 Number of inactive employees in 2013: 0

Number of active employees in 2014: 1881  
 Number of inactive employees in 2014: 410

```
In [12]: averagel=(2291+1881)/2
turnover_rate_13_14= 410/averagel
turnover_rate_13_14
```

```
Out[12]: 0.1965484180249281
```

```
In [13]: turnover_year=org.groupby("cutoff_date")["turnover"].mean()
print(turnover_year)

cutoff_date
2014-12-31    0.178961
Name: turnover, dtype: float64
```

```
In [14]: # Number of new joinees every 3 years
```

```
In [15]: joining_1994_1997 = org[(org["date_of_joining"]>="1994-01-01") &
                                   (org["date_of_joining"]<="1997-12-31")]

joining_1997_2000 = org[(org["date_of_joining"]>="1997-01-01") &
                           (org["date_of_joining"]<="2000-12-31")]

joining_2000_2003 = org[(org["date_of_joining"]>="2000-01-01") &
                           (org["date_of_joining"]<="2003-12-31")]

joining_2003_2006= org[(org["date_of_joining"]>="2003-01-01") &
                           (org["date_of_joining"]<="2006-12-31")]

joining_2006_2009 = org[(org["date_of_joining"]>="2006-01-01") &
                           (org["date_of_joining"]<="2009-12-31")]

joining_2009_2012 = org[(org["date_of_joining"]>="2009-01-01") &
                           (org["date_of_joining"]<="2012-12-31")]


joining_2012_2015 = org[(org["date_of_joining"]>="2012-01-01") &
                           (org["date_of_joining"]<="2015-12-31")]
```

## Showing turnover rate level-wise

```
In [16]: turnover_level=org.groupby("level")["turnover"].mean()  
print(turnover_level)  
sns.catplot(kind="bar",  
            x="level",  
            y="turnover",  
            data=org,  
            ci=None)  
plt.xticks(rotation=90)  
plt.show()
```

*""Turnover rate is highest at Analyst and Specialist level""*


```
level  
Analyst                0.215087  
Assistant Manager      0.036458  
Director               0.000000  
Manager               0.043478  
Senior Manager        0.000000  
Specialist            0.148571  
Vice President        0.000000  
Name: turnover, dtype: float64
```



Out[16]: 'Turnover rate is highest at Analyst and Specialist level'

```
In [17]: turnover_level.plot()  
plt.xticks(rotation=90)
```

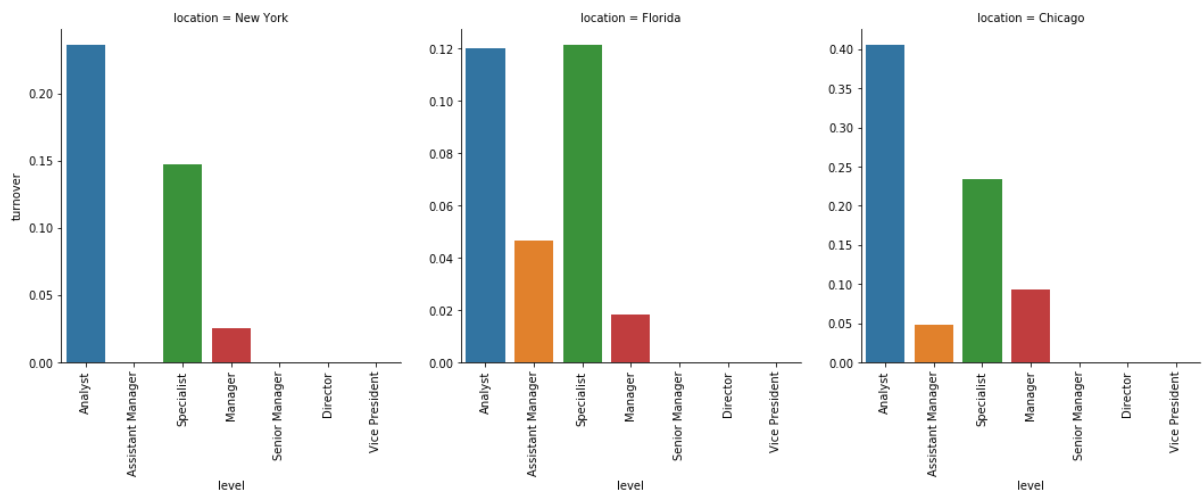
```
Out[17]: (array([0., 1., 2., 3., 4., 5., 6.]), <a list of 7 Text xticklabel objects>)
```





```
In [18]: g=sns.catplot(data=org,
                      kind="bar",
                      x="level",
                      y="turnover",
                      col="location",
                      col_wrap=3,
                      ci=None,
                      sharey=False)
for ax in g.axes:
    plt.setp(ax.get_xticklabels(), visible= True, rotation=90)
#plt.xticks(rotation=90)
plt.show()

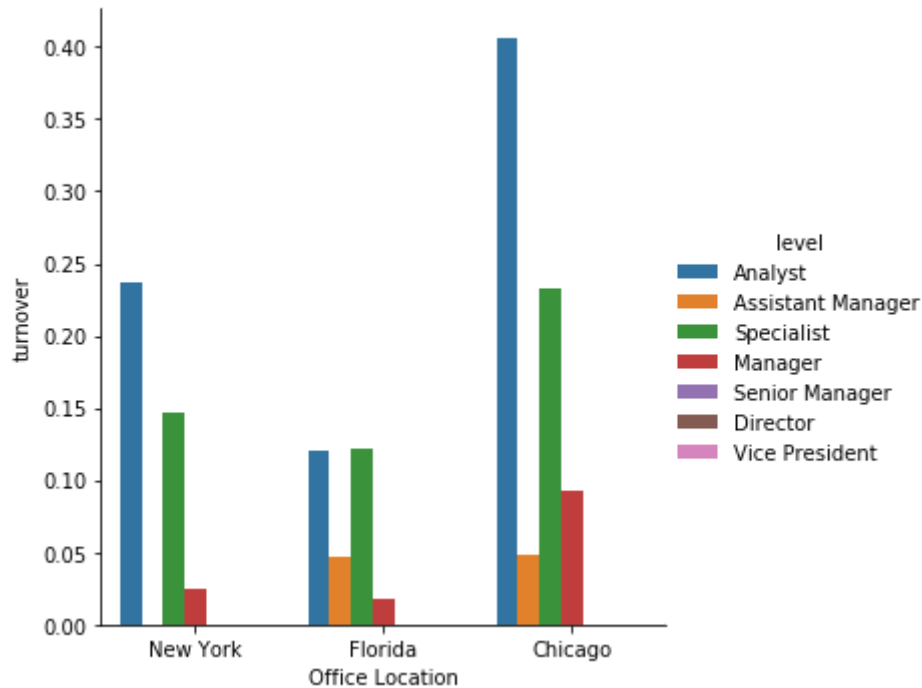
"""Turnover rate is highest at Analyst and Specialist level at all locations."""
```



```
Out[18]: 'Turnover rate is highest at Analyst and Specialist level at all locations.'
```

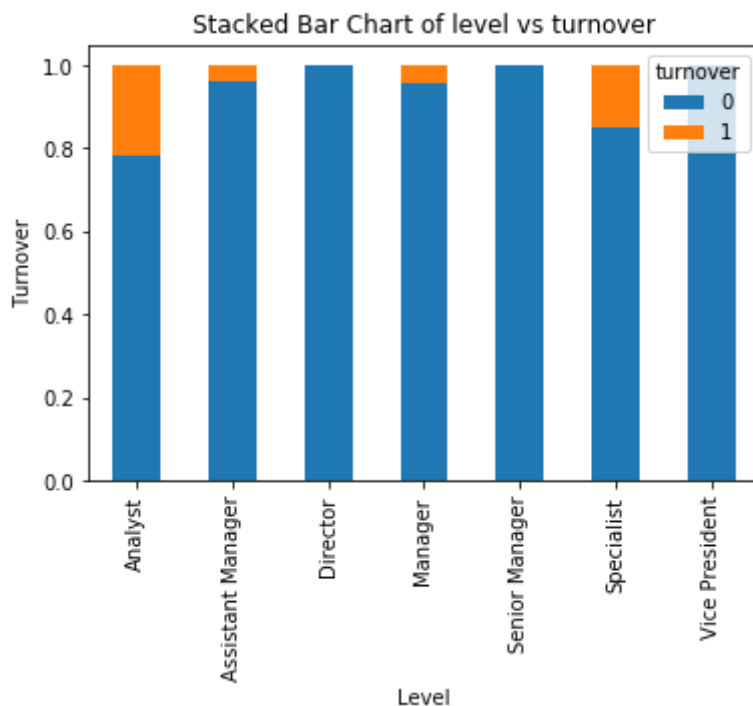
```
In [19]: ans= org.groupby("location")["turnover"].agg(np.mean)
print(ans)
sns.catplot(kind="bar", x="location", y="turnover", hue="level", data=org, ci=None)
plt.xlabel("Office Location")
plt.show()
"""Turnover rate is high in Chicago as compared to other locations"""
```

```
location
Chicago      0.325641
Florida      0.105513
New York     0.202591
Name: turnover, dtype: float64
```



```
Out[19]: 'Turnover rate is high in Chicago as compared to other locations'
```

```
In [20]: table=pd.crosstab(org.level,org.turnover)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of level vs turnover')
plt.xlabel('Level')
plt.ylabel('Turnover')
plt.savefig('Level vs Turnover')
```



We can calculate categorical means for other categorical variables such as education and marital status to get a more detailed sense of our data.

```
In [21]: org.groupby("location").mean()
```

Out[21]:

	turnover	emp_age
location		
Chicago	0.325641	31.247436
Florida	0.105513	31.011692
New York	0.202591	27.915548

```
In [22]: org.groupby("level").mean()  
# the following gives same output too  
#org.pivot_table(values=["turnover","emp_age"], index="level", aggfunc  
="mean")
```

Out[22]:

	turnover	emp_age
level		
Analyst	0.215087	28.480611
Assistant Manager	0.036458	34.354167
Director	0.000000	39.000000
Manager	0.043478	36.166667
Senior Manager	0.000000	39.600000
Specialist	0.148571	31.326286
Vice President	0.000000	40.000000

```
In [23]: org.groupby("gender").mean()
```

Out[23]:

	turnover	emp_age
gender		
Female	0.175389	29.091231
Male	0.180556	30.267424

## BINS CREATION

```
In [24]: # making bins of Employee age
#org["emp_age"].min()---22
#org["emp_age"].max()---58
bins=[20,30,40,50,60]
bin_labels=["20-30","30-40","40-50","50-60"]
org["age_range"]=pd.cut(org["emp_age"],bins=bins,right=True, labels=bin_labels)
org
```

Out[24]:

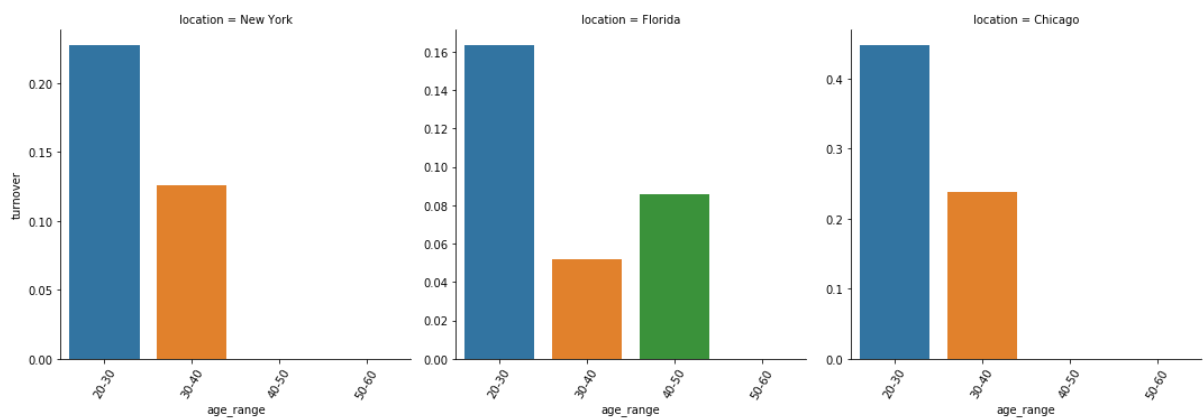
	emp_id	status	turnover	location	level	date_of_joining	date_of_birth	last_working_d
0	E11061	Inactive	1	New York	Analyst	2012-03-22	1992-03-22	2014-11
1	E1031	Inactive	1	New York	Analyst	2012-09-03	1992-10-01	2014-05
2	E6213	Inactive	1	New York	Analyst	2012-06-01	1992-06-02	2014-04
3	E5900	Inactive	1	New York	Analyst	2012-03-22	1991-12-19	2014-05
4	E3044	Inactive	1	Florida	Analyst	2012-03-29	1991-10-12	2014-01
...	...	...	...	...	...	...	...	...
2286	E8787	Inactive	1	Florida	Assistant Manager	2013-06-03	1965-07-20	2014-02
2287	E12351	Active	0	Florida	Assistant Manager	1996-09-10	1965-12-06	2021-03
2288	E6282	Active	0	New York	Manager	2012-08-23	1960-05-18	2021-03
2289	E5195	Active	0	Florida	Analyst	2002-07-01	1972-07-01	2021-03
2290	E11037	Active	0	Florida	Manager	1994-01-21	1956-10-15	2021-03

2291 rows × 15 columns

```
In [25]: g=sns.catplot(data=org,
                      kind="bar",
                      x="age_range",
                      y="turnover",
                      col="location",
                      col_wrap=3,
                      ci=None,
                      sharey=False)
for ax in g.axes:
    plt.setp(ax.get_xticklabels(), visible= True, rotation=60)

plt.show()
```

*""Turnover is highest among 20-30 years of age range.  
However, in Florida, about 8% of turnover is among employees between 40-50 years of age. ""*



```
Out[25]: 'Turnover is highest among 20-30 years of age range.\nHowever, in Florida, about 8% of turnover is among employees between 40-50 years of age.'
```

## Observations

Approximately 18% of employees left the organization from 2013-2014 year.

The turnover is highest at Analyst and Specialist level.

Turnover rate is high in Chicago as compared to other locations.

The mean age of employees at Analyst and Specialist level is about 29 and 31 respectively.

There is not much difference in turnover rate across male and female employees. The gender does not seem a strong predictor for the turnover.

```
In [26]: org.groupby("level")["location"].value_counts()
```

```
Out[26]: level          location
Analyst      Florida      682
            New York      656
            Chicago      266
Assistant Manager Florida    129
            New York      42
            Chicago      21
Director     Florida      1
Manager      Florida      55
            Chicago      43
            New York      40
Senior Manager Florida      3
            New York      2
Specialist   Florida     181
            New York     109
            Chicago      60
Vice President Florida      1
Name: location, dtype: int64
```

**Talent segments** We will exclude top level management and middle management from analysis. We will focus our analysis on Analysts and Specialists as they form majority of the workforce. See the count below:

```
In [27]: org["level"].value_counts()
```

```
Out[27]: Analyst      1604
Specialist      350
Assistant Manager  192
Manager         138
Senior Manager    5
Vice President    1
Director          1
Name: level, dtype: int64
```

```
In [28]: entry_level=["Analyst", "Specialist"]
org2 =org[org.level.isin(entry_level)]
```

```
In [29]: #org1=org[(org["level"]=="Analyst") | (org["level"]=="Specialist")]
#org1
```

```
In [30]: turnover_level=org.groupby("generation")["turnover"].mean()
print(turnover_level)
```

```
generation
Baby Boomers    0.000000
Generation X     0.086047
Millennials     0.200646
Name: turnover, dtype: float64
```

## Exploratory Data Analysis results:

# Combining data from different HR sources

```
In [31]: survey=pd.read_csv("survey.csv")
survey.head()
```

Out[31]:

	mgr_id	mgr_effectiveness	career_satisfaction	perf_satisfaction	work_satisfaction
0	E1003	0.76	0.76	0.71	0.82
1	E10072	0.65	0.67	0.56	0.84
2	E10081	0.80	0.82	0.73	0.84
3	E10234	0.65	0.63	0.75	0.70
4	E1026	0.70	1.00	1.00	0.92

```
In [32]: survey.shape
```

Out[32]: (350, 5)

```
In [33]: #left_df.merge(right_df, on='user_id', how='left', indicator=True)
org_2=org.merge(survey,on="mgr_id", how='left', indicator=False)
org_2.head()
```

Out[33]:

	emp_id	status	turnover	location	level	date_of_joining	date_of_birth	last_working_date
0	E11061	Inactive	1	New York	Analyst	2012-03-22	1992-03-22	2014-11-09
1	E1031	Inactive	1	New York	Analyst	2012-09-03	1992-10-01	2014-05-06
2	E6213	Inactive	1	New York	Analyst	2012-06-01	1992-06-02	2014-04-30
3	E5900	Inactive	1	New York	Analyst	2012-03-22	1991-12-19	2014-09-04
4	E3044	Inactive	1	Florida	Analyst	2012-03-29	1991-10-12	2014-01-23

```
In [34]: org_2.shape
```

Out[34]: (2291, 19)



```
In [35]: rating=pd.read_csv("rating.csv")
rating.head()
```

Out[35]:

	emp_id	rating
0	E8	Acceptable
1	E9	Acceptable
2	E12	Acceptable
3	E15	Acceptable
4	E34	Acceptable

```
In [36]: org_3=org_2.merge(rating, on="emp_id", how='left', indicator=False)
org_3.head()
```

Out[36]:

	emp_id	status	turnover	location	level	date_of_joining	date_of_birth	last_working_date
0	E11061	Inactive	1	New York	Analyst	2012-03-22	1992-03-22	2014-11-09
1	E1031	Inactive	1	New York	Analyst	2012-09-03	1992-10-01	2014-05-06
2	E6213	Inactive	1	New York	Analyst	2012-06-01	1992-06-02	2014-04-30
3	E5900	Inactive	1	New York	Analyst	2012-03-22	1991-12-19	2014-09-04
4	E3044	Inactive	1	Florida	Analyst	2012-03-29	1991-10-12	2014-01-23

```
In [37]: org_3.shape
```

Out[37]: (2291, 20)

```
In [38]: #org_3.info()
```

## Feature Engineering---using a different dataset

Process of using domain knowledge to create new variables which help you discover new insights. Crucial step before building a predictive model.

### New features added:

**Age difference**= manager age- employee age Views, handling pressure, Expectation, Work ethics Younger employees feel Older employees cannot deal with changing work pace.

**JOB hopper**- Job hopper is a person who switches job frequently for financial or career advancement opportunities. Recruiters and hiring managers views job hoppers in a negative light. There is a high chance they would quit soon. Job hop index = Total work experience/Number of companies worked.

### Employee Tenure

Inactive employees: date\_of\_joining & last\_working\_date

Active employee : date\_joining & cutoff\_date(study period end date)

The more the number of years an employee works in the organization, it is less likely that he/she will quit.


```
In [39]: org_final=pd.read_csv("org_final.csv", parse_dates=["date_of_joining", "last_working_date", "cutoff_date"])
org_final.shape
```

```
Out[39]: (1954, 34)
```

```
In [40]: #org_final.info()
```

```
In [41]: mgr_eff=org_final.groupby("status")["mgr_effectiveness"].agg(np.mean)
print(mgr_eff)
sns.catplot(kind="box",
            data=org_3,
            x="status",
            y="mgr_effectiveness")
plt.show()
```

```
status
Active      0.708169
Inactive    0.568516
Name: mgr_effectiveness, dtype: float64
```



## Observations:

Mean manager effectiveness seemed to be higher in active employees as compared to inactive employees.


Age difference between manager and employee seem to be higher in inactive employees than active. Therefore could be a predictor for turnover.

Median job hop index for active and inactive employee are similar.

Age difference between manager and employee

```
In [42]: sns.set_context("paper")
org_final["age_diff"] = org_final["mgr_age"] - org_final["emp_age"]
age_2 = org_final.groupby("status")["age_diff"].mean()
print(age_2)
g = sns.catplot(kind="bar",
                x="status",
                y="age_diff",
                data=org_final)
g.fig.suptitle("Manager and Employee's age difference vs Turnover ")
g.set(xlabel="Status", ylabel="Age difference b/w employee and manager")
plt.show()
```

```
status
Active      4.691946
Inactive    5.982393
Name: age_diff, dtype: float64
```



## Job- Hop Index

If its high means person didnt switch much. If its low, means employee switched a lot.

```
In [43]: # replacing infinite values(inf) with zero
org_final["job_hop_index"] = org_final["total_experience"].div(org_final[
    "no_previous_companies_worked"]).replace(np.inf, 0)
org_final[["job_hop_index"]]
```

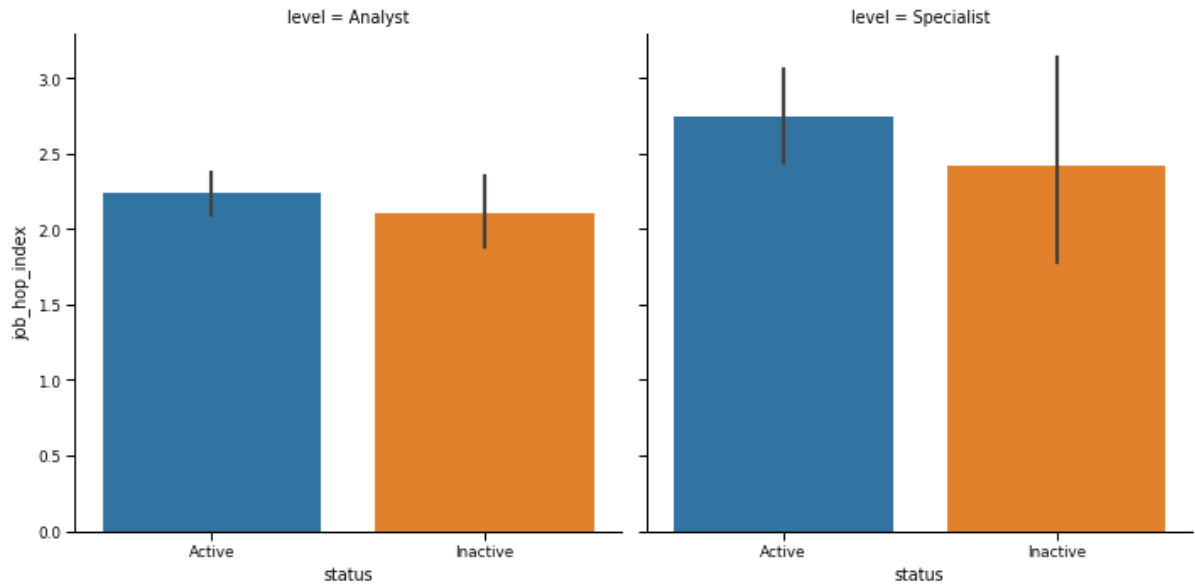
Out[43]:

	job_hop_index
0	0.000000
1	0.542222
2	2.850000
3	0.952000
4	0.000000
...	...
1949	4.050000
1950	2.600000
1951	2.425000
1952	1.428571
1953	0.920000

1954 rows × 1 columns

```
In [44]: sns.catplot(kind="bar",
                    y="job_hop_index",
                    x="status",
                    data=org_final,
                    col="level")
plt.show()
```

*""Analysts have a lower mean job hop index- means they are more likely to leave than Specialists""*



Out[44]: 'Analysts have a lower mean job hop index- means they are more likely to leave than Specialists'

```
In [45]: org_final["date_of_joining"].fillna(0)
org_final["last_working_date"].fillna(0)
org_final["date_of_joining"]=pd.to_datetime(org_final["date_of_joining"]
])
org_final["last_working_date"]=pd.to_datetime(org_final["last_working_date"]
])
org_final["cutoff_date"]=pd.to_datetime(org_final["cutoff_date"])
#org_final.info()
```

## Calculating tenure in years


```
In [46]: # Different calculations for 'Active' and "Inactive" employees
org_final.loc[org_final["status"]=="Inactive", "tenure"] = (org_final["last_working_date"]-org_final["date_of_joining"])/np.timedelta64(1, 'Y')
org_final.loc[org_final["status"]=="Active", "tenure"] = (org_final["cutoff_date"]-org_final["date_of_joining"])/np.timedelta64(1, 'Y')
#org_final
```

```
In [47]: ten_lev= org_final.groupby("level")["tenure"].agg(np.median)
print(ten_lev)
```

```
level
Analyst      3.438811
Specialist   5.411473
Name: tenure, dtype: float64
```

```
In [48]: sns.catplot(data=org_final, x="status", y="tenure", kind="box")
plt.show()
```

*""The median tenure of inactive employees is less than the tenure of active employees.""*




```
Out[48]: 'The median tenure of inactive employees is less than the tenure of active employees.'
```

## Viewing distribution of compensation among all employees

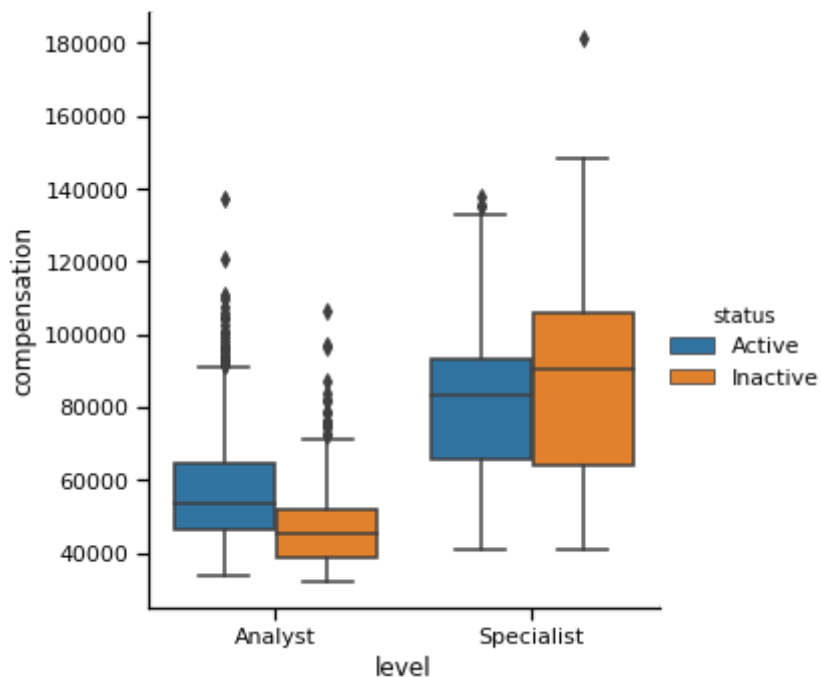
```
In [49]: sns.set_context("notebook")
plt.hist(x="compensation", data=org_final, bins=40)
plt.xticks(rotation=90)
plt.plot()
```

Out[49]: []



The distribution is right skewed.

```
In [50]: sns.set_context("notebook")
sns.catplot(kind="box",
            y="compensation",
            x="level",
            data=org_final,
            hue="status")
plt.show()
```





**Compensation varies across jobs and levels. At Analyst level, median compensation of Inactive employees is low when compared to Active employees.(Compensation could be a reason for leaving).**

**At specialist level, the median compensation of inactive employees was higher than active employees(some other reason for leaving).**

Employees always expect that they are paid fairly compared to their co-workers, and hence, maintaining internal pay parity is important.

Competitiveness of each employee's pay can be assessed by Compa-ratio. In this exercise, you'll derive compa-ratio as:

Compa-ratio= Actual compensation/Median compensation

Median compensation is used by organizations to estimate the typical pay for any role/level. This metric helps the organization to correct the compensation of employees who are way below the median compensation.

Remember, median is also known as the 50th percentile. Exactly 50 percent of people make less than the median and 50 percent make more.

```
In [51]: median_comp=org_final.groupby("level")["compensation"].agg(np.median)
         median_comp
```

```
Out[51]: level
Analyst      51840
Specialist    83496
Name: compensation, dtype: int64
```


```
In [52]: org_final.loc[org_final["level"]=="Analyst", "compa_ratio"]=org_final["compensation"]/51840
org_final.loc[org_final["level"]=="Specialist", "compa_ratio"]=org_final["compensation"]/83496
org_final
```

Out[52]:

	emp_id	status	location	level	gender	emp_age	rating	mgr_rating	mgr_reporter
0	E10012	Active	New York	Analyst	Female	25.09	Above Average	Acceptable	
1	E10025	Active	Chicago	Analyst	Female	25.98	Acceptable	Excellent	
2	E10027	Active	Orlando	Specialist	Female	33.40	Acceptable	Above Average	
3	E10048	Active	Chicago	Specialist	Male	24.55	Acceptable	Acceptable	
4	E10060	Active	Orlando	Analyst	Male	31.23	Acceptable	Acceptable	
...	...	...	...	...	...	...	...	...	
1949	E9960	Active	Orlando	Analyst	Male	27.81	Excellent	Acceptable	
1950	E9977	Active	Orlando	Analyst	Male	27.64	Above Average	Above Average	
1951	E9980	Active	New York	Specialist	Male	27.63	Acceptable	Acceptable	
1952	E9992	Active	Chicago	Specialist	Male	28.34	Acceptable	Acceptable	
1953	E9993	Active	Orlando	Analyst	Female	26.25	Acceptable	Acceptable	

1954 rows × 38 columns

```
In [53]: sns.catplot(kind="box",  
                    x="level",  
                    y="compa_ratio",  
                    data=org_final)  
plt.show()
```

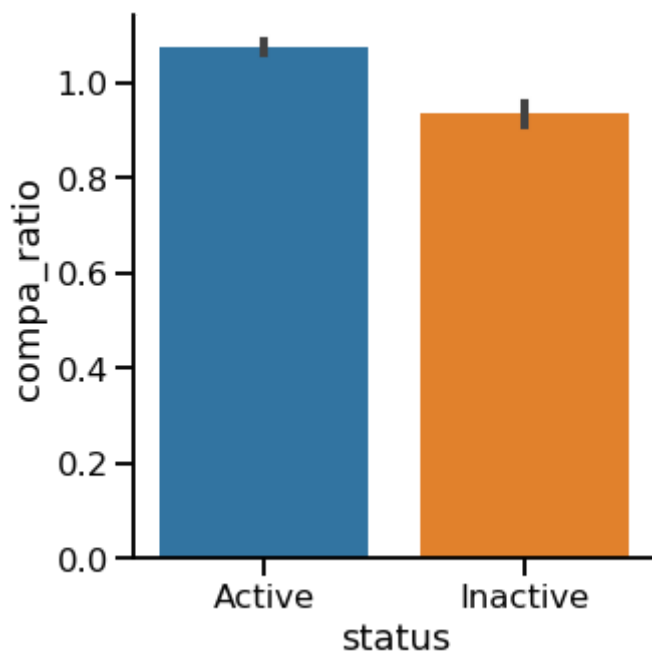


**Median compa-ratio is same in both analysts and Specialists. At specialist level, median is closer to 75th percentile, more number of employees have compa-ratio less than the median.**

```
In [54]: sns.set_context("talk")
sns.catplot(data=org_final,
            x="status",
            y="compa_ratio",
            kind="bar")

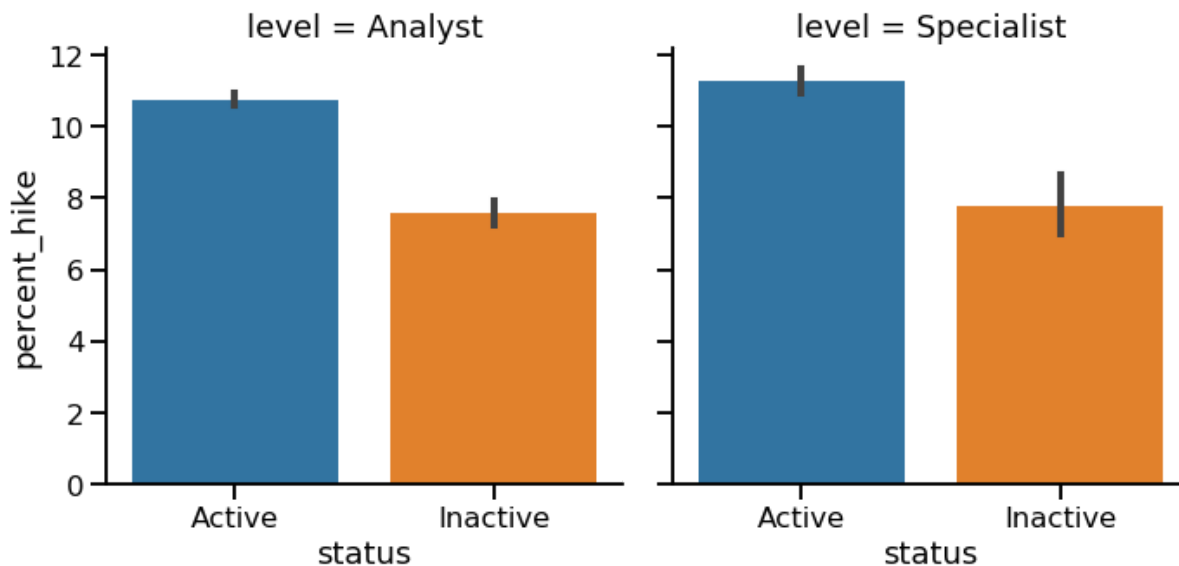
plt.show()

"""A greater proportion of inactive employees were paid less than median
compensation."""
```




```
Out[54]: 'A greater proportion of inactive employees were paid less than median
compensation.'
```

```
In [55]: sns.set_context("talk")
sns.catplot(data=org_final,
            x="status",
            y="percent_hike",
            kind="bar",
            col="level")
plt.show()
```



```
In [56]: sns.set_context("talk")
sns.catplot(data=org_final,
            x="turnover",
            y="percent_hike",
            kind="bar",
            hue="level",
            col="location",
            ci=None)
plt.xticks(rotation=90)
plt.show()
```



## Predictor variable's contribution

```
In [57]: # only for quantitative variables
corr_matrix_f = org_final.corr().round(2)
corr_matrix_f
```

Out[57]:


	emp_age	mgr_reportees	mgr_age	mgr_tenure	compensation	per
emp_age	1.00	-0.10	0.12	0.02	0.61	
mgr_reportees	-0.10	1.00	-0.14	0.03	-0.14	
mgr_age	0.12	-0.14	1.00	0.20	0.13	
mgr_tenure	0.02	0.03	0.20	1.00	-0.04	
compensation	0.61	-0.14	0.13	-0.04	1.00	
percent_hike	0.01	-0.08	0.04	-0.01	0.15	
hiring_score	0.02	0.01	0.02	0.02	0.07	
no_previous_companies_worked	0.05	-0.03	-0.01	-0.00	0.02	
distance_from_home	-0.08	0.12	-0.02	-0.01	-0.11	
total_dependents	0.18	0.06	-0.03	-0.02	0.10	
no_leaves_taken	-0.07	0.10	0.02	-0.02	-0.09	
total_experience	0.86	-0.07	0.09	0.01	0.52	
monthly_overtime_hrs	-0.03	0.03	-0.02	-0.01	-0.01	
turnover	-0.17	0.23	-0.03	-0.05	-0.18	
mgr_effectiveness	0.06	-0.21	-0.05	-0.05	0.08	
career_satisfaction	0.04	-0.02	-0.05	0.04	0.04	
perf_satisfaction	-0.02	-0.03	-0.08	-0.07	-0.03	
work_satisfaction	-0.05	-0.04	-0.08	-0.02	-0.06	
age_diff	-0.61	-0.05	0.71	0.15	-0.33	
job_hop_index	0.30	-0.01	0.02	-0.02	0.19	
tenure	0.59	-0.08	0.08	0.09	0.44	
compa_ratio	0.52	-0.09	0.10	-0.04	0.76	

22 rows × 22 columns

```
In [58]: # distance_from_home, compensation, mgr_effectiveness, perf_satisfaction,
age_diff, tenure,
```

```
In [59]: plt.figure(figsize=(10,10))
sns.heatmap(corr_matrix_f)
```

```
Out[59]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa725bd45d0>
```



Weight of evidence (WOE) and Information value (IV) are simple, yet powerful techniques to perform variable transformation and selection. These concepts have huge connection with the logistic regression modeling technique. It is widely used in credit scoring to measure the separation of good vs bad customers.

```
In [60]: # removing variables which have no predictive power

# emp_id, mgr_id ----are id columns
# date_of_joining, last_working_date, cutoff_date ----tenure is a linear
combination of these columns.
# mgr_age, emp_age-----age_diff is a linear combination of these column
s.
# department-----
# status ---is same as turnover
```

```
In [61]: from xverse.transformer import MonotonicBinning
independent_var=['location', 'level', 'gender', 'emp_age', 'rating', 'mgr_
_rating', 'mgr_reportees', 'mgr_age',
                'mgr_tenure', 'compensation', 'percent_hike', 'hiring_sc
ore', 'hiring_source',
                'no_previous_companies_worked', 'distance_from_home', 'to
tal_dependents', 'marital_status', 'education',
                'promotion_last_2_years', 'no_leaves_taken', 'total_exp
erience',
                'monthly_overtime_hrs', 'department', 'turnover', 'mgr_ef
fectiveness',
                'career_satisfaction', 'perf_satisfaction', 'work_satis
faction', 'age_diff', 'job_hop_index', 'tenure', 'compa_ratio']
x_check=org_final[independent_var]
y_check=org_final["turnover"]
clf = MonotonicBinning()
clf.fit(x_check, y_check)

#print(clf.bins)
```

```
Out[61]: MonotonicBinning(cardinality_cutoff=5, custom_binning=None, feature_nam
es='all',
                           force_bins=4, max_bins=20, prefix=None)
```



```
In [62]: from xverse.transformer import WOE
```

```
clf = WOE()
clf.fit(x_check, y_check)
```

```
print(clf.woe_df.head()) #Weight of Evidence transformation dataset
```

	Variable_Name	Category	Count	Event	Non_Event	Event
_Rate \						
0	age_diff	(-15.071, 2.48]	652	100	552	0.1
53374						
1	age_diff	(2.48, 7.59]	651	144	507	0.2
21198						
2	age_diff	(7.59, 24.98]	651	153	498	0.2
35023						
3	career_satisfaction	(-0.001, 0.75]	689	151	538	0.2
19158						
4	career_satisfaction	(0.75, 0.87]	631	128	503	0.2
02853						

	Non_Event_Rate	Event_Distribution	Non_Event_Distribution	WOE
\				
0	0.846626	0.251889	0.354528	-0.341798
1	0.778802	0.362720	0.325626	0.107882
2	0.764977	0.385390	0.319846	0.186418
3	0.780842	0.380353	0.345536	0.096001
4	0.797147	0.322418	0.323057	-0.001980

	Information_Value
0	0.051302
1	0.051302
2	0.051302
3	0.007064
4	0.007064

```
/opt/anaconda3/lib/python3.7/site-packages/pandas/core/series.py:853: RuntimeWarning: divide by zero encountered in log
  result = getattr(ufunc, method)(*inputs, **kwargs)
```

```
In [63]: print(clf.iv_df) #Information value dataset

# Less than 0.02 - Not useful for prediction
# 0.02 - 0.1 - Weak predictive power
# 0.1 - 0.3 - Medium predictive power
# 0.3 - 0.5 - Strong predictive power
# >0.5 - Suspicious predictive power.
```

	Variable_Name	Information_Value
28	total_dependents	1.081841
5	distance_from_home	0.878127
23	percent_hike	0.769505
16	mgr_effectiveness	0.586315
27	tenure	0.472047
3	compensation	0.413557
21	no_leaves_taken	0.392402
26	rating	0.386937
13	location	0.296302
18	mgr_reportees	0.287404
2	compa_ratio	0.285904
20	monthly_overtime_hrs	0.181906
7	emp_age	0.174745
6	education	0.125387
29	total_experience	0.110376
25	promotion_last_2_years	0.099799
31	work_satisfaction	0.063526
0	age_diff	0.051302
11	job_hop_index	0.044011
19	mgr_tenure	0.039950
24	perf_satisfaction	0.032243
15	mgr_age	0.030725
12	level	0.027265
9	hiring_score	0.026808
14	marital_status	0.025881
17	mgr_rating	0.021722
10	hiring_source	0.008774
1	career_satisfaction	0.007064
22	no_previous_companies_worked	0.001342
8	gender	0.000040
4	department	0.000000
30	turnover	0.000000

## Predictive Analysis

### Logistic Regression

Its a classification problem where we are trying to predict whether turnover will occur or not. Our dependent variable(turnover) is categorical.

# Checking multicollinearity

Inappropriate use of dummies causes multicollinearity too. I removed the categorical data to check for multicollinearity

```
In [64]: non_cat= {"tenure", "compensation", "no_leaves_taken", "compa_ratio", "hi
ring_score", "percent_hike",
                 "emp_age", "distance_from_home", "mgr_reportees", "monthly_
overtime_hrs",
                 "total_experience"}
x_df=pd.get_dummies(org_final[non_cat])
```

```
In [65]: from statsmodels.stats.outliers_influence import variance_inflation_fact
or
# For each X, calculate VIF and save in dataframe
vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_df.values, i) for i in
range(x_df.shape[1])]
vif["features"] = x_df.columns
vif.round(1)
```

Out[65]:

	VIF Factor	features
0	8.0	tenure
1	127.5	hiring_score
2	10.8	percent_hike
3	6.4	distance_from_home
4	228.2	emp_age
5	31.3	compensation
6	39.2	compa_ratio
7	7.6	mgr_reportees
8	3.4	monthly_overtime_hrs
9	23.5	total_experience
10	4.1	no_leaves_taken

## Removing employee age

```
In [66]: x_df=x_df.drop(columns="emp_age")
```

```
In [67]: from statsmodels.stats.outliers_influence import variance_inflation_factor
or
# For each X, calculate VIF and save in dataframe
vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_df.values, i) for i in
range(x_df.shape[1])]
vif["features"] = x_df.columns
vif.round(1)
```

Out[67]:

	VIF Factor	features
0	7.1	tenure
1	44.8	hiring_score
2	10.8	percent_hike
3	6.4	distance_from_home
4	30.8	compensation
5	38.4	compa_ratio
6	7.6	mgr_reportees
7	3.4	monthly_overtime_hrs
8	11.3	total_experience
9	4.1	no_leaves_taken

## Removing compa ratio

```
In [68]: x_df=x_df.drop(columns="compa_ratio")
```

```
In [69]: vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_df.values, i) for i in
range(x_df.shape[1])]
vif["features"] = x_df.columns
vif.round(1)
```

Out[69]:

	VIF Factor	features
0	6.8	tenure
1	40.3	hiring_score
2	10.8	percent_hike
3	6.3	distance_from_home
4	16.3	compensation
5	7.6	mgr_reportees
6	3.4	monthly_overtime_hrs
7	10.9	total_experience
8	4.1	no_leaves_taken

## Removing hiring score

```
In [70]: x_df=x_df.drop(columns="hiring_score")
```

```
In [71]: vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_df.values, i) for i in
range(x_df.shape[1])]
vif["features"] = x_df.columns
vif.round(1)
```

Out[71]:

	VIF Factor	features
0	6.8	tenure
1	7.6	percent_hike
2	5.5	distance_from_home
3	14.7	compensation
4	6.3	mgr_reportees
5	3.3	monthly_overtime_hrs
6	10.6	total_experience
7	3.8	no_leaves_taken

## Removing compensation

```
In [72]: x_df=x_df.drop(columns="compensation")
```

```
In [73]: vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(x_df.values, i) for i in
range(x_df.shape[1])]
vif["features"] = x_df.columns
vif.round(1)
```

Out[73]:

	VIF Factor	features
0	6.3	tenure
1	6.5	percent_hike
2	5.4	distance_from_home
3	6.3	mgr_reportees
4	3.3	monthly_overtime_hrs
5	8.3	total_experience
6	3.8	no_leaves_taken

## CHOSEN PREDICTORS AND DUMMY CODING

```
In [74]: # predictors chosen with very little multicollinearity and categorical v  
         # ariables added back  
         predictors={"mgr_effectiveness","tenure", "no_leaves_taken","rating","em  
         p_age", "location", "distance_from_home" ,"mgr_reportees", "monthly_over  
         time_hrs", "promotion_last_2_years",  
         "total_experience","percent_hike"}  
         target="turnover"  
         x=pd.get_dummies(org_final[predictors], drop_first=True)  
         y=org_final[target]  
         #x.head()
```

## Using statsmodels to display summary

```
In [75]: import statsmodels.api as sm  
logit_model=sm.Logit(y, x)  
result=logit_model.fit()  
print(result.summary())
```

Optimization terminated successfully.

Current function value: 0.184229

Iterations 9

### Logit Regression Results

```

=====
=====
Dep. Variable:          turnover    No. Observations:
1954
Model:                  Logit       Df Residuals:
1938
Method:                 MLE        Df Model:
15
Date:                   Thu, 15 Apr 2021    Pseudo R-squ.:
0.6350
Time:                   00:27:41    Log-Likelihood:
-359.98
converged:              True        LL-Null:
-986.32
Covariance Type:        nonrobust    LLR p-value:              7.8
86e-258
=====
=====

```

		coef	std err	z	P> z
[0.025	0.975]				
-----					
tenure		-0.3937	0.075	-5.266	0.000
-0.540	-0.247				
mgr_effectiveness		-4.4025	0.561	-7.848	0.000
-5.502	-3.303				
percent_hike		-0.4570	0.047	-9.634	0.000
-0.550	-0.364				
distance_from_home		0.2192	0.016	13.417	0.000
0.187	0.251				
emp_age		-0.0632	0.034	-1.843	0.065
-0.130	0.004				
mgr_reportees		0.1104	0.019	5.914	0.000
0.074	0.147				
monthly_overtime_hrs		0.2015	0.027	7.481	0.000
0.149	0.254				
total_experience		0.0070	0.050	0.140	0.889
-0.092	0.106				
no_leaves_taken		0.1233	0.013	9.332	0.000
0.097	0.149				
location_New York		1.3033	0.277	4.699	0.000
0.760	1.847				
location_Orlando		-0.4246	0.230	-1.843	0.065
-0.876	0.027				
rating_Acceptable		-0.1233	0.242	-0.509	0.611
-0.598	0.352				
rating_Below Average		-1.5228	0.436	-3.490	0.000
-2.378	-0.668				
rating_Excellent		-0.6544	0.630	-1.038	0.299
-1.890	0.581				
rating_Unacceptable		-2.6284	0.766	-3.430	0.001
-4.130	-1.127				
promotion_last_2_years_Yes		-0.0173	0.286	-0.061	0.952



-0.577            0.542

=====

=====

\*\*I noticed several variables are insignificant. In multiple regression models, this can happen due to multicollinearity.

```
In [76]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=1)
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

```
(1367, 16)
(1367,)
(587, 16)
(587,)
```

## Logistic Regression

```
In [77]: # Regress all variables in the dataframe on target variable 'turnover'
model = LogisticRegression(solver='liblinear', random_state=1)
model.fit(x_train, y_train)
```

```
Out[77]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2',
                             random_state=1, solver='liblinear', tol=0.0001, verbose=0,
                             warm_start=False)
```

### Predicting

```
In [78]: y_pred=model.predict(x_test)
y_pred_proba=model.predict_proba(x_test)

result=pd.DataFrame({"Observed": y_test,
                    "Predicted": y_pred,
                    "p(0)": [p[0] for p in y_pred_proba],
                    "p(1)": [p[1] for p in y_pred_proba]})

print(result)
print(y_pred.shape)
```

	Observed	Predicted	p(0)	p(1)
1812	0	0	0.992061	0.007939
1576	0	0	0.989677	0.010323
1926	0	0	0.986105	0.013895
1378	0	0	0.996522	0.003478
1446	0	0	0.990683	0.009317
...	...	...	...	...
110	0	0	0.998302	0.001698
1105	0	0	0.938261	0.061739
1481	0	0	0.982100	0.017900
1247	0	0	0.996325	0.003675
955	0	1	0.480526	0.519474

```
[587 rows x 4 columns]
(587,)
```

## Model Evaluation

```
In [79]: prediction_train = model.predict(x_train)
prediction_valid = model.predict(x_test)
print("Accuracy of model on training set is:", accuracy_score(y_train, prediction_train))
print("Accuracy of model on test set is:", accuracy_score(y_test, prediction_valid))
```

```
Accuracy of model on training set is: 0.9180687637161667
Accuracy of model on test set is: 0.9369676320272572
```

```
In [80]: #Confusion Matrix
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[470  11]
 [ 26  80]]
```

The result is telling us that we have 467+72 correct predictions and 14+34 incorrect predictions.

```
In [81]: # Classification report
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.98	0.96	481
1	0.88	0.75	0.81	106
accuracy			0.94	587
macro avg	0.91	0.87	0.89	587
weighted avg	0.94	0.94	0.94	587

```
In [82]: score_ = accuracy_score(y_test, y_pred)
prec_score = precision_score(y_test, y_pred)
f1=f1_score(y_test,y_pred)

print("The accuracy score for the test is:", score_, "\n")
print(f1)
print(regressionSummary(y_test,y_pred))
print(model.score(x_test,y_test))
```

The accuracy score for the test is: 0.9369676320272572

0.8121827411167513

Regression statistics

Mean Error (ME) : 0.0256  
 Root Mean Squared Error (RMSE) : 0.2511  
 Mean Absolute Error (MAE) : 0.0630  
 None  
 0.9369676320272572

The precision is the ratio  $tp / (tp + fp)$  where  $tp$  is the number of true positives and  $fp$  the number of false positives. The precision is intuitively the ability of the classifier to not label a sample as positive if it is negative.

The recall is the ratio  $tp / (tp + fn)$  where  $tp$  is the number of true positives and  $fn$  the number of false negatives. The recall is intuitively the ability of the classifier to find all the positive samples.

The F-beta score can be interpreted as a weighted harmonic mean of the precision and recall, where an F-beta score reaches its best value at 1 and worst score at 0.\*\*

The F-beta score weights the recall more than the precision by a factor of beta. beta = 1.0 means recall and precision are equally important.

The support is the number of occurrences of each class in  $y_{test}$ .

## Decision trees

```
In [83]: # i had to remove all categorical variables from predictors
```

```
In [84]: # predictors chosen with very little multicollinearity and categorical v
          ariables added back
predictors_tree={"mgr_effectiveness","tenure", "no_leaves_taken","emp_ag
e", "distance_from_home" ,"mgr_reportees", "monthly_overtime_hrs",
               "total_experience","percent_hike"}

target="turnover"
x=pd.get_dummies(org_final[predictors_tree], drop_first=True)
y=org_final[target]
#x.head()
x_train, x_test, y_train,y_test=train_test_split(x,y,test_size=0.3, rand
om_state=1)
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

(1367, 9)
(1367,)
(587, 9)
(587,)
```

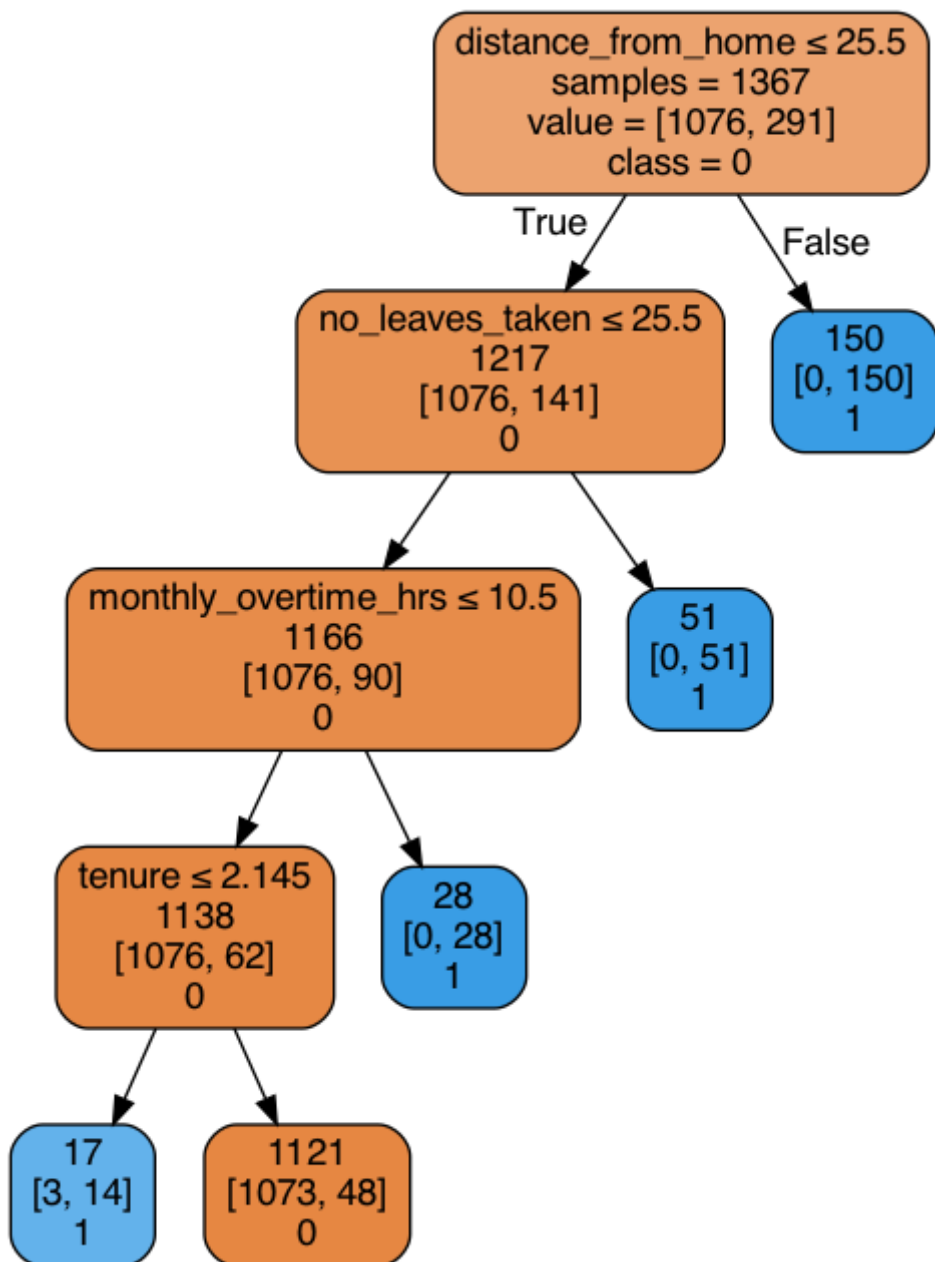
```
In [85]: #Build decision tree
          from sklearn.tree import DecisionTreeClassifier
          dt = DecisionTreeClassifier(max_depth = 4)
          dt.fit(x_train, y_train)
```

```
Out[85]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
4,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=No
ne,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

```
In [86]: from dmba import plotDecisionTree, gainsChart, liftChart
from dmba import classificationSummary, regressionSummary

plotDecisionTree(dt, feature_names=org_final[predictors_tree].columns, class_names=dt.classes_)
```

Out[86]:



```
In [87]: importances = dt.feature_importances_
imp_features = pd.DataFrame({'feature': x_train.columns, 'importance': importances})
imp_features = imp_features.sort_values('importance', ascending = False)
print(imp_features)
```

	feature	importance
3	distance_from_home	0.577893
8	no_leaves_taken	0.230354
6	monthly_overtime_hrs	0.135248
0	tenure	0.056504
1	mgr_effectiveness	0.000000
2	percent_hike	0.000000
4	emp_age	0.000000
5	mgr_reportees	0.000000
7	total_experience	0.000000

```
In [88]: #get the accuracy score
from sklearn.metrics import accuracy_score
# Get the prediction for both train and test
prediction_train = dt.predict(x_train)
prediction_valid = dt.predict(x_test)

# Measure the accuracy of the model for both train and test sets
print("Accuracy on train is:", accuracy_score(y_train, prediction_train))
print("Accuracy on test is:", accuracy_score(y_test, prediction_valid ))
```

```
Accuracy on train is: 0.9626920263350403
Accuracy on test is: 0.9761499148211243
```

## Retention Strategy

In [89]: *# filter data for only active employees*

```
org_final_active=org_final[org_final["status"]=="Active"]
org_final_active
```

Out[89]:

	emp_id	status	location	level	gender	emp_age	rating	mgr_rating	mgr_reporter
0	E10012	Active	New York	Analyst	Female	25.09	Above Average	Acceptable	
1	E10025	Active	Chicago	Analyst	Female	25.98	Acceptable	Excellent	
2	E10027	Active	Orlando	Specialist	Female	33.40	Acceptable	Above Average	
3	E10048	Active	Chicago	Specialist	Male	24.55	Acceptable	Acceptable	
4	E10060	Active	Orlando	Analyst	Male	31.23	Acceptable	Acceptable	
...	...	...	...	...	...	...	...	...	
1949	E9960	Active	Orlando	Analyst	Male	27.81	Excellent	Acceptable	
1950	E9977	Active	Orlando	Analyst	Male	27.64	Above Average	Above Average	
1951	E9980	Active	New York	Specialist	Male	27.63	Acceptable	Acceptable	
1952	E9992	Active	Chicago	Specialist	Male	28.34	Acceptable	Acceptable	
1953	E9993	Active	Orlando	Analyst	Female	26.25	Acceptable	Acceptable	

1557 rows × 38 columns

In [90]: *# predictors chosen with very little multicollinearity and categorical variables added back*

```
predictors=["mgr_effectiveness","tenure", "no_leaves_taken","rating","emp_age", "location", "distance_from_home", "mgr_reportees", "monthly_over_time_hrs", "promotion_last_2_years", "total_experience", "percent_hike"]
target="turnover"
x_active=pd.get_dummies(org_final_active[predictors], drop_first=True)
y_active=org_final_active[target]
```

```
In [91]: # Directly asking the model to predict for all employees who are active
# No test and training sets
y_pred=model.predict(x_active)
y_pred_proba=model.predict_proba(x_active)

result=pd.DataFrame({"Predicted": y_pred,
                    "p(0)": [p[0] for p in y_pred_proba],
                    "p(1)": [p[1] for p in y_pred_proba]})

print(result)
```

	Predicted	p(0)	p(1)
0	0	0.996922	3.077946e-03
1	0	0.999999	1.131291e-06
2	0	1.000000	1.155836e-13
3	0	0.999916	8.384654e-05
4	0	0.999997	2.971293e-06
...	...	...	...
1552	0	0.999556	4.442536e-04
1553	0	1.000000	1.976737e-07
1554	0	1.000000	1.723229e-07
1555	0	0.999544	4.559267e-04
1556	0	0.995317	4.683232e-03

[1557 rows x 3 columns]

```
In [92]: org_final_active["risk_probability"]=result["p(1)"]
org_final_active["risk_probability"]=org_final_active["risk_probability"]
.fillna(0)
#org_final_active
```

0 < Employees with turnover probability < 0.5 --- NO risk bucket

0.5 < Employees with turnover probability < 0.6 ---- Low risk

0.6 < Employees with turnover probability < 0.8 --- High risk bucket

Employees with turnover probability > 0.8 --- High risk bucket

```
In [93]: df=org_final_active.copy()
```

```
In [94]: # Finding high risk employees(employee ids) most likely to leave

high_risk_df= df[df["risk_probability"]>=0.8]
medium_risk_df=df[(df["risk_probability"]<0.8)&(df["risk_probability"]>=
0.6)]
low_risk_df= df[(df["risk_probability"]>=0.5)& (df["risk_probability"]<
0.6)]
no_risk_df= df[df["risk_probability"]<0.5]

high_risk_df["risk_probability"].count()
```

Out[94]: 3



```
In [95]: medium_risk_df["risk_probability"].count()
```

```
Out[95]: 5
```

```
In [96]: low_risk_df["risk_probability"].count()
```

```
Out[96]: 5
```

```
In [97]: no_risk_df["risk_probability"].count()
```

```
Out[97]: 1544
```

## Retention Strategy

### High risk

If a high risk employee is a high performer and has high potential, immediate action planning needed.

A. Engage in a conversation with this employee to generally understand the perspective about work and future plans.

B. Ask the employee's manager to have a conversation and explore the engagement levels and concerns, if any.

### Medium risk

A. Medium term action planning.

B. Have one-on-one or open house discussion.

C. Keep tracking of any behavioral change.

### Low risk

A. Long-term action planning

B. Keep tracking for any behavioral change.

C. Have open house discussion.

### No risk

No action required.

Cost of employee turnover:

Costs to off-board employee

Cost-per-hire for replacement

Transition costs, including opportunity costs

**The most important predictor- percent\_hike. Consider giving a better percentage hike.**

```
In [98]: org_fin=pd.read_csv("org_final.csv")
org_fin.fillna(0)
```

Out[98]:

	emp_id	status	location	level	gender	emp_age	rating	mgr_rating	mgr_reporter
0	E10012	Active	New York	Analyst	Female	25.09	Above Average	Acceptable	
1	E10025	Active	Chicago	Analyst	Female	25.98	Acceptable	Excellent	
2	E10027	Active	Orlando	Specialist	Female	33.40	Acceptable	Above Average	
3	E10048	Active	Chicago	Specialist	Male	24.55	Acceptable	Acceptable	
4	E10060	Active	Orlando	Analyst	Male	31.23	Acceptable	Acceptable	
...	...	...	...	...	...	...	...	...	
1949	E9960	Active	Orlando	Analyst	Male	27.81	Excellent	Acceptable	
1950	E9977	Active	Orlando	Analyst	Male	27.64	Above Average	Above Average	
1951	E9980	Active	New York	Specialist	Male	27.63	Acceptable	Acceptable	
1952	E9992	Active	Chicago	Specialist	Male	28.34	Acceptable	Acceptable	
1953	E9993	Active	Orlando	Analyst	Female	26.25	Acceptable	Acceptable	

1954 rows × 34 columns

```
In [99]: org_hike_count=org_fin["percent_hike"].value_counts()
org_hike_count
```

```
Out[99]: 9      229
10     220
11     214
12     198
14     181
13     180
7      180
8      170
15     100
6       67
5       49
3       41
4       30
0       30
16      16
19      16
17      15
18      12
2        6
Name: percent_hike, dtype: int64
```

```
In [100]: #org_final["percent_hike"] = org_final["percent_hike"].astype('category')

cat_1=org_fin.loc[(org_fin["percent_hike"] >=0 ) & (org_fin["percent_hike"] < 5 ), "percent_hike"]
cat_2=org_fin.loc[(org_fin["percent_hike"] >=5 ) & (org_fin["percent_hike"] < 10 ), "percent_hike"]
cat_3=org_fin.loc[(org_fin["percent_hike"] >=10 ) & (org_fin["percent_hike"] < 15 ), "percent_hike"]
print(cat_2)
```

```
1      8
3      8
5      8
7      9
8      9
..
1943   9
1944   9
1945   8
1952   7
1953   8
Name: percent_hike, Length: 695, dtype: int64
```